

Top 10 Key Takeaways

Chunks 1–9 · 13:24–14:53 · Instructors Madeleine Woods (MW) and Marlon Kuzmick (MK), with the executive-director demo from Becca

Distilled from the diarized session transcript at [day-1-recap/20260608-transcript.md](#). Faculty questions from Jonah, Ian, Christine, Jordan, and Jungyoon, among others.

- "Strings in, strings of numbers out" — that is all an LLM is.** The day's most repeated mantra. The Tiktokenizer demo on `unhappily` (split into `un` + `app` + `ily`) makes it visceral: language gets chopped into chunks "no human decided," reduced to integers, and the model predicts more integers. A "very inhuman sort of machine" that mimics patterns of human knowledge — and everything downstream (memory, the system prompt, skills, even `/compact`) is just more text.
- LLMs recognize patterns but cannot actually calculate — so make them call a tool.** The five-digit × five-digit demo: first three digits, last three digits, and order of magnitude come out right; the middle is hallucinated. Marlon: it's "close to truth, capable of persuading you that they are true if you're not careful." Ask for a Python function instead — "we're really leveraging what it's trained to do well." LLM-alone-vs.-LLM-plus-tool recurs everywhere.
- It will be confidently — and sycophantically — wrong.** The live *Coriolanus* close-reading demo: Claude insists Act 1 Scene 9 doesn't exist, then ("I apologize. I'm so sorry. You're totally right") flips and analyzes it once the scene is pasted in. Hallucination plus sycophancy on stage — the naïve "do my homework" failure mode faculty most need to internalize.
- Don't ask for an *image* of a data viz — ask for an interactive artifact and attach the CSV.** The population-pyramid contrast: a diffusion model produces something pretty but wrong (pixels, not data); Claude artifacts open code tools, read the CSV's structure, and generate an actual interactive. As Marlon put it for the recipe-website moment later: "you don't need to know code at all, you just need to know to ask for an HTML doc."
- The "light at the end of the tunnel" is Becca's demos — what faculty can actually build.** A Monty Hall lesson auto-assembled from existing course materials into the new LXP (video clip → quiz → vibe-coded interactive → reveal clip). Handwriting grading on in-class assessments, where Claude reads the student's diagonal cross-outs and `UUURRR` lattice notation, and gets it wrong "at a lesser rate than my TFs." A blackboard-video lecture turned into course notes.
- The context window = Memento (or "a brilliant student who hasn't done the reading").** Every new chat is amnesiac. What loads first and invisibly: the **system prompt** (Anthropic publishes theirs — safety plus style; "you can't edit it"), then **memory** (also just text — a file you can toggle on). Then prompts and responses pile in until you hit the wall: ~200K tokens for older models, 1M for newer. Marlon's alternative for non-Memento-watchers: don't let the brilliant student fill the window with smart-sounding stuff that isn't grounded.
- Primacy, recency, and "context rot" — the middle of a long thread quietly loses weight.** The first tokens (system prompt) and the most recent instruction pull the most weight; everything in between degrades as the window grows. The move: refresh aggressively — start new chats, run `/compact` before you have to, and shift the ratio toward your own vetted, high-signal context so Claude isn't dominating with its own long outputs.
- Cowork's killer feature is local-file access — and its main risk is yes-fatigue.** Chat (cloud-only) can't touch your files; Cowork and Code can read and write them. Point Cowork at one specific folder (subfolders are included — "I wouldn't give it access to my desktop or, you know, my root drive"). Watch the permission prompts: the Homer-Simpson-at-the-reactor warning — Claude will keep asking until you stop reading and just click yes. Create a dedicated `DeveLopment/` folder.
- The recipe-cards demo is the whole workshop in five minutes.** Point Cowork at a folder of weirdly-named recipe photos; ask it to rename each by what's on the card, write a markdown doc of the recipe, add historical context, and finally produce "one single stylish HTML file" embedding all of them. In one prompt: batch file ops + image reading + transcription + creative writing + HTML output. Translate the same move to syllabi, manuscripts, marginalia, blackboard photos, a teaching journal.
- The recipe mental model — ingredients (context) + instructions (prompt) → dish (output).** The day's closing analogy and the take-home homework: list the ingredients you'd hand a "brilliant student who hasn't done the reading" (syllabus, readings, P-sets, course videos), specify the output precisely (a new file? slides plus a lecture script?), and write the prompts that connect them. Skills, scripts, and tools — coming on later days — are just ways of packaging the same recipe so you stop copy-pasting.

SECONDARY POINTS WORTH KEEPING

- **Model tiers.** Opus 4.8 (most capable, slowest, priciest); Sonnet 4.6 (workhorse); Haiku (cheapest). Dial reasoning effort high → low for speed during demos.
- **Markdown is the lingua franca** for feeding context — also HTML and other structured formats. Matthew Schwartz's online interactive physics textbook is the scale-up example.
- **"Coding is just reading and writing files."** Said in passing while introducing GitHub — but it's the frame that lets non-coders work in Cowork and Code without flinching.
- **"You are X" is the speaker-audience inversion.** With an LLM you're at "interaction zero" — the LLM is the speaker, you set the role and the form. ChatGPT's original "You are a helpful assistant" is the ur-version.
- **Chain prompts, and draft prompts with another LLM.** Translation courses use this: produce a draft, ask "what are the top five problems with this," then fix. Madeleine grows a "super prompt" inside one chat, then starts a fresh chat with it.
- **Compaction sands off rough edges.** As Claude summarizes its own context to keep going, your original vision drifts — worth flagging for

- **"Ten top tips from yesterday" is the model for your own course materials.** Hand students the curated repo and let Claude be their index — grounding their learning in context you've vetted.
- **3D / spatial layout is still weak.** Madeleine still hand-edits XY coordinates in vibe-coded frontends.
- **/compact is editable.** Madeleine wrote her own **/mw-compact** because she disagrees with some of Anthropic's defaults.
- **Sycophancy is a feature of training, not a bug** — flag for faculty so they don't read "you're absolutely right" as confirmation.